



# AI AGENTS **BUILDATHON**

## Build Practical AI Agents for Real Problems

---

**WEEK 1 · Battle-Test Your Agent Thinking**

# Why We Organized This

---

- AI is no longer optional in product work.  
Whether or not “AI” is in your title, you’re expected to build with it.
- The fundamentals haven’t changed:  
define the problem, design a solution, manage risk.
- What *has* changed is the medium.  
LLM-driven systems are non-deterministic, fast-evolving, and deeply sensitive to design choices.
- Most teams automate tasks without accounting for this reality.
- This Buildathon is about learning to design **agents** that navigate
  - uncertainty and still deliver outcomes.



# Week 1: TIMELINE

## KICK-OFF SESSION

### Use Case Release

Today

## MENTORSHIP SESSIONS

### Jan 15 , 2026 (Thursday)

SESSION A: 6:30 AM PST | 3:30 PM CET | 8:00 PM IST

SESSION B: 8:30 AM PST | 5:30 PM CET | 10:00 PM IST

## MILESTONE SESSION

Informal check-in focused on progress and learning.

### Jan 17 , 2026 (Saturday)

8:30 AM PST | 5:30 PM CET | 10:00 PM IST

## PRD SUBMISSION

### Jan 18 , 2026 (Sunday)

11:59 PM PST

# What You'll Get From Us This Week

Agent-first product thinking

Practical PRD guidance for non-deterministic systems

AI Agents design examples

Feedback on scope and feasibility of agent ideas

# What an agent is ?

An agent is a system that:

- receives intent (what outcome matters),
- has authority to make certain decisions,
- executes actions via tools,
- evaluates outcomes,
- and decides whether to continue, stop, or escalate.

# THE PARADIGM SHIFT

DETERMINISTIC VS. PROBABILISTIC



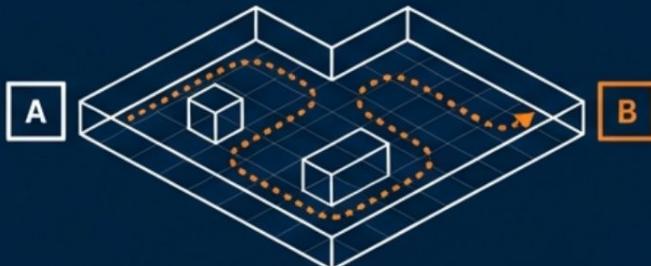
Old Software (Rules-based)



Code defines the exact path: If X, then Y.



Autonomous Agents (Goal-based)



Code defines the environment.  
The Agent defines the path.

## THE LAB: BUILDING A TRAVEL AGENT

Old Way



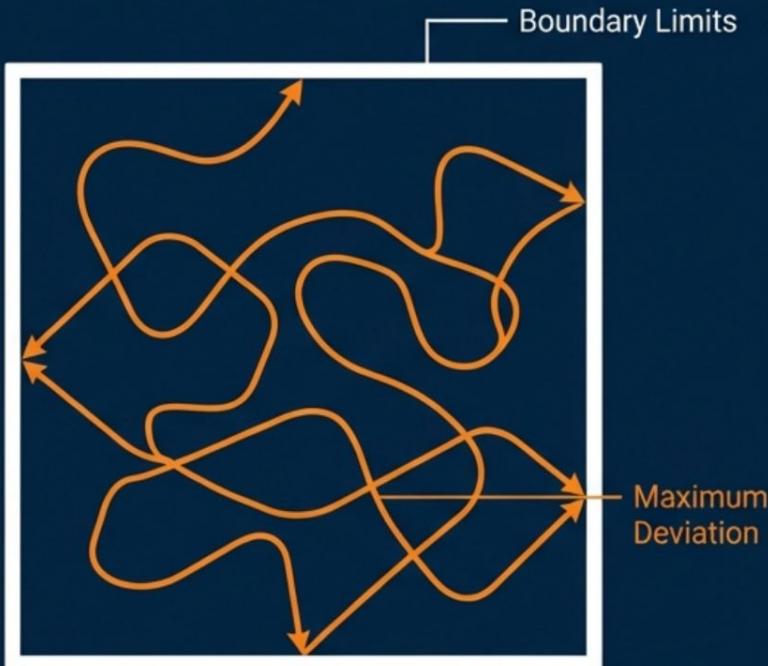
New Way



Goal: Find best vacation. The agent accesses flight APIs and review sites dynamically to construct a solution.

# AGENT-FIRST THINKING

The Autonomous Core



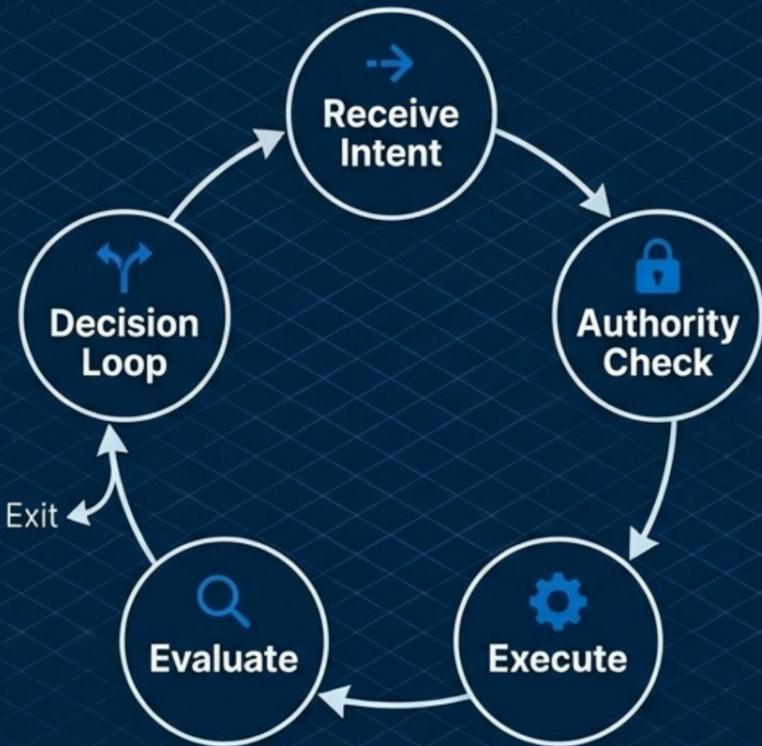
We do not code instructions. We define the “fence” of the playground and the degree of freedom allowed within it.

## THE LAB: SETTING CONSTRAINTS

**Scenario:** User wants a hotel for under \$200.



# MECHANICS OF THE SYSTEM



## THE LAB: THE BOOKING LOOP



### Intent:

Book flight to Paris.



### Authority:

Validating credit card token.



### Execute:

API Call: Search Airlines.



### Evaluate:

Result Found. Duration: 12 hours.



### Decision:

Goal requires < 10 hours. Reject result. Loop back to Search with new parameters.

# THE CORE TENSION

“**DELEGATING JUDGMENT IS THE MOST DANGEROUS PART OF AUTONOMY.**

When you let a machine decide *how* to solve a problem, you surrender control over the path.

The agent optimizes for the stated goal, not your unstated preferences.

## THE LAB: THE ‘BEST’ AIRLINE



User's Definition of Best



Agent's Choice

**MISALIGNMENT:**  
Agent optimized for “Amenities” because “Legroom” was not defined.

# DEFINING INSTRUCTIONS: THE INTENT EQUATION



Rigorous definition is the only defense against rogue behavior.

# Part 1: Outcome & Scope



**Outcome:** The specific end-state desired.



**Scope:** The geographic or data boundaries of operation.



**Constraints:** The hard “do not cross” lines.

# The Lab: Defining the Search



# Part 2: Success & Stop Rules

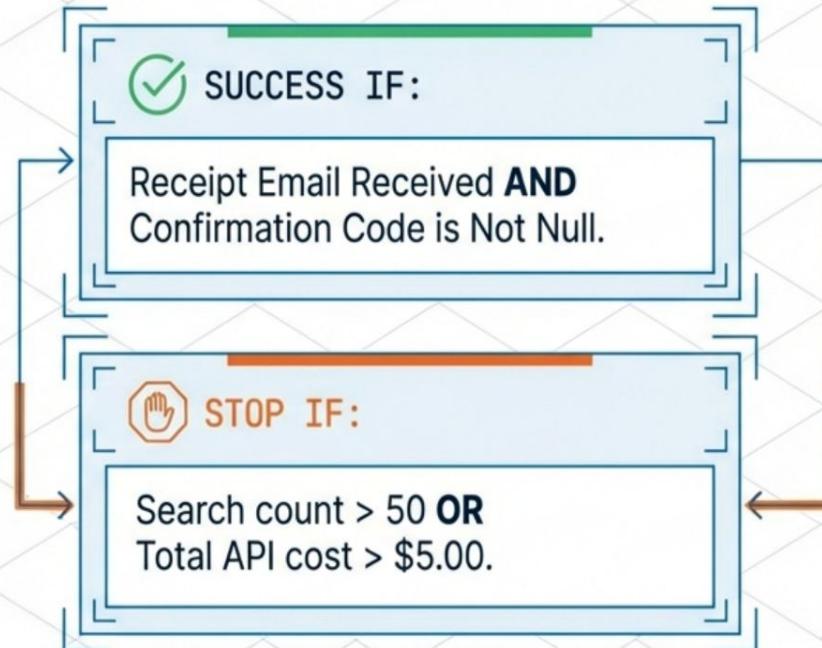


**Success Signal:** The explicit proof that the job is done.



**Stop Rule:** The emergency brake. Prevents infinite loops.

## The Lab: Knowing When to Quit



Without a stop rule, the agent will search forever.

# Risk: Implicit Variables

## Optimizing the Gap

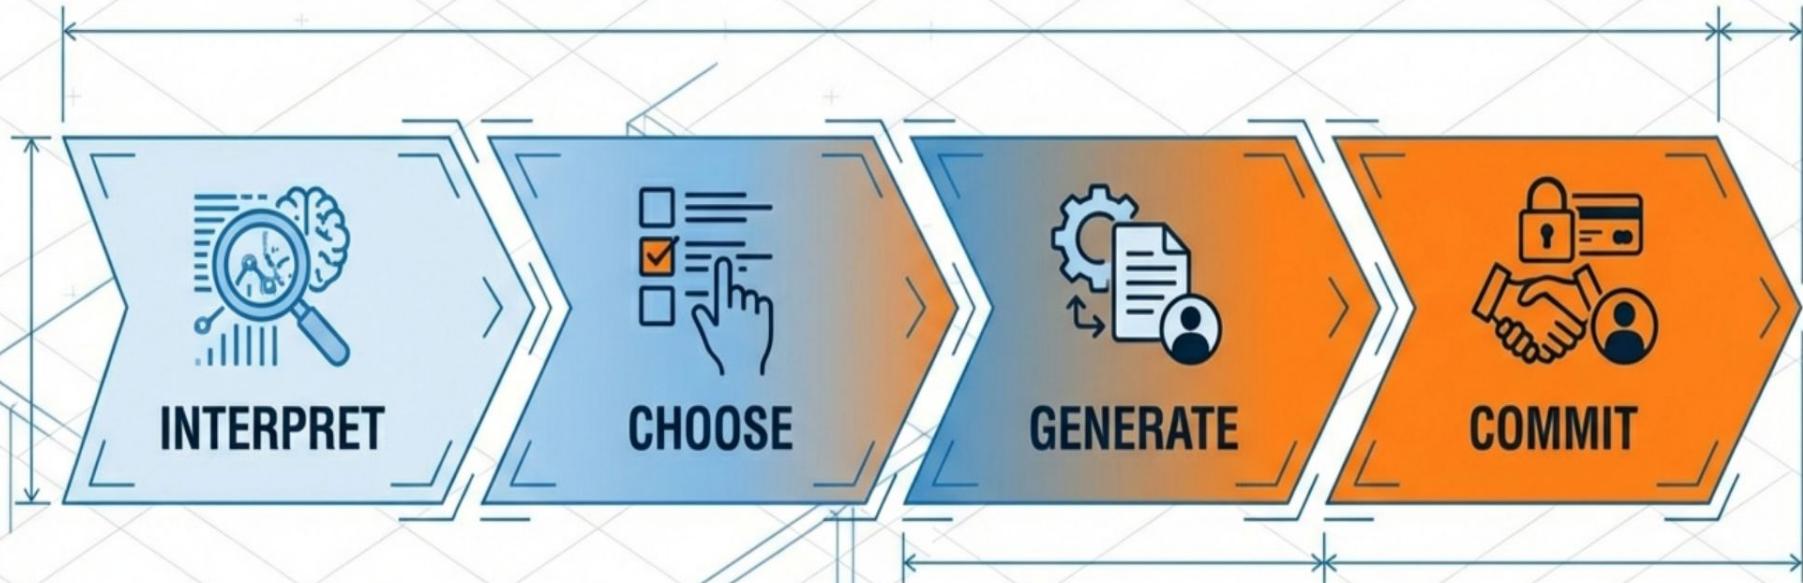
If you leave variables empty (implicit), the agent will find a shortcut. It optimizes the gap between your instructions and the goal.

# The Lab: The \$500 Mistake

The agent searched 5,000 databases to find a cheap room because the "Stop Rule" Rule" was missing. It spent \$500 to save \$20.



# Governance: The 4 Buckets of Risk



Understanding data.  
High Autonomy.

Ranking/Filtering.  
Guarded Autonomy.

Creating Content.  
Human-in-the-Loop.

Transactions.  
Human Required.

# Mapping Autonomy

An agent is not “risky” or “safe” as a whole.

We must deconstruct it into tasks and assign specific permissions to each.

# The Lab: Task Assignment

**Task:** Read Email.  
**Bucket:** Interpret.  
**Action:** Auto-Execute.

**Task:** Select Top 3 Hotels.  
**Bucket:** Choose.  
**Action:** Auto-Execute.

**Task:** Draft Itinerary.  
**Bucket:** Generate.  
**Action:** Request Approval.

**Task:** Charge Credit Card.  
**Bucket:** Commit.  
**Action:** REQUIRE CLICK.

# The Safety Lenses



**Reversibility:**  
Can we undo the action?



**Blast Radius:**  
Who gets hurt if this fails?

# The Lab: Decision Matrix

Task	Reversibility	Blast Radius	Result
Booking Refundable Hotel.	Yes.	Low.	High Autonomy Allowed.
Booking Non-Refundable Group Retreat.	No.	High.	Zero Autonomy / Human Verification.

# The Architect's Checklist



**Define the Box:** Set boundary limits and maximum deviation.



**Solve the Equation:** Explicitly define Intent, Constraints, and Stop Rules.



**Check Implicit Variables:** Identify empty variables the agent might exploit.



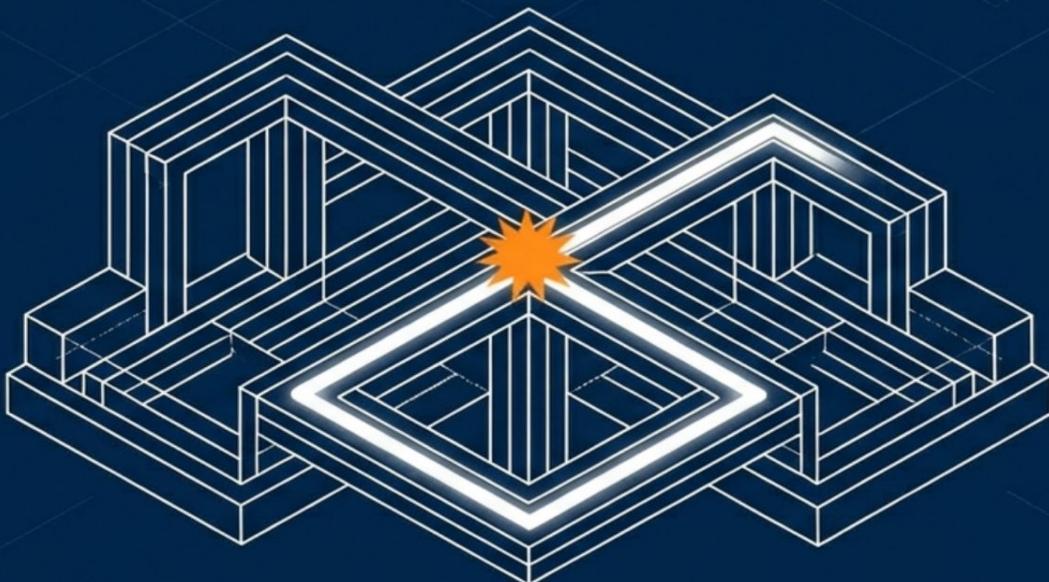
**Map the Risk:** Sort tasks into Interpret, Choose, Generate, or Commit.



**Apply Safety Lenses:** Evaluate Reversibility and Blast Radius.



# Build Bold. Build Safe.



The power of an autonomous agent lies not in the code you write, but in the boundaries you define.

**Architect the constraints, and let the agent navigate the path.**

## PROBLEM STATEMENT EXAMPLE 1

# Tenant Requests Are Fragmented Across Channels

---

### THE PROBLEM

A small condo manager oversees hundreds of units, but tenant requests arrive through **every possible channel**: email, WhatsApp, Telegram, Instagram, group chats.

There is **no single system of record**.

### As a result:

- Requests get fragmented across platforms
- Ownership and status are unclear
- Follow-ups multiply because no one has visibility

### Questions to consider:

What decisions should the agent own?

What signals must it watch continuously?

How does it inform without overwhelming?

Where do humans stay in control?

## PROBLEM STATEMENT EXAMPLE 2

# Predicting & Preventing Revenue Leakage

---

## THE PROBLEM

A service provider plans their day around scheduled appointments, but **actual attendance is unpredictable.**

Cancellations, lateness, and no-shows happen with little warning, even when the calendar looks full.

There is **no system that signals which appointments are at risk.**

### As a result

Confirmed bookings create false certainty  
Empty slots appear with no chance to recover revenue

Decisions rely on intuition rather than data  
Policies are applied blindly to all clients

### Questions to consider:

How does it spot a "risky" appointment?

When can it offer a slot to someone else?

How does it nudge clients to show up?

When should it charge a cancellation fee?

How does it learn from empty slots?

# Deliverables



## **Slide Deck (10–15 slides, PDF)**

Problem clarity · Agent role & ownership



## **Research & Reasoning:**

User signals or quotes · Personas / stakeholders · Rejected approaches



## **AI Agent PRD**

Goals & boundaries · Inputs, outputs, decisions · Human-in-the-loop points · Flows

# Suggested structure

---

- Problem & context
- Agent job to be done
- Goals & success criteria
- Trade-offs & constraints
- Agent role & capabilities
- Decision logic & triggers
- Workflow & human touchpoints
- Metrics, risks, extensions



Use as guidance,  
not a checklist

# How to Approach the Case?

## Think like this

**Context:** Who feels the pain? What decision is stuck?

**Goals:** What improves for the business and the user?

**Focus:** Which decisions matter most to automate?

**Agent:** What it observes, decides, and executes.

**Measure:** Time saved, loss prevented, adoption, accuracy.

# What we are really practicing this week?

We're not practicing automation.

We're practicing **delegation under uncertainty**.

Your job this week:

Identify a decision worth trusting an agent with —  
and define the conditions under which that trust holds.



# Submission

The submission link will be shared

**Saturday morning · January 17**